

The aim of this data analysis is twofold: first, to develop a model that can accurately forecast property prices in Beijing, China, and secondly, to construct a model that categorizes properties as either underpriced or overpriced relative to their original listing price.

1. Initially, the focus is on cleaning and preprocessing the data, followed by constructing a regression model capable of predicting property prices based on various property features.
2. Moving on to the second phase of the analysis, K-means clustering is utilized to group properties according to their features. Subsequently, each cluster is examined to identify properties that significantly deviate from the cluster's average price, thereby identifying underpriced and overpriced properties.

```
In [1]: # Loading the necessary Libraries
import re
import folium
import numpy as np
import pandas as pd
import seaborn as sns
import lightgbm as lgb
from gower import gower_matrix
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, silhouette_score, r2_score
```

Data Dictionary:

1. url: the url which fetches the data
2. id: the id of transaction
3. Lng: Longitudes
4. Lat: Latitudes
5. Cid: community id
6. tradeTime: the time of transaction
7. DOM: active days on market
8. followers: the number of people following the transaction.
9. totalPrice: the total price
10. price: the average price by square
11. square: the total area of the house
12. livingRoom: the number of living rooms
13. drawingRoom: the number of drawing rooms
14. kitchen: the number of kitchen
15. bathroom: the number of bathrooms
16. floor: the height of the house
17. buildingType: including tower(1), bungalow(2), combination of plate and tower(3), plate(4).
18. constructionTime: the time of construction
19. renovationCondition: including other(1), rough(2),Simplicity(3), hardcover(4)

20. buildingStructure: including unknow(1), mixed(2), brick and wood(3), brick and concrete(4), steel(5) and steel-concrete composite (6).
21. ladderRatio: the proportion between number of residents on the same floor and number of elevator of ladder. It describes how many ladders a resident have on average.
22. elevator have (1) or not have elevator(0)
23. fiveYearsProperty: if the owner have the property for less than 5 years


In [2]: `# Setting the seed and Loading the dataset`
`np.random.seed(99)`
`df = pd.read_csv(r'C:\Users\himan\OneDrive\Desktop\Beijing_Housing.csv', encoding='utf-8')`

C:\Users\himan\AppData\Local\Temp\ipykernel_20144\1512002212.py:3: DtypeWarning: Columns (1,11,12,14) have mixed types. Specify dtype option on import or set low_memory=False.
`df = pd.read_csv(r'C:\Users\himan\OneDrive\Desktop\Beijing_Housing.csv', encoding='ISO-8859-1')`

Out[2]:

	url	id	Lng	Lat	
0	https://bj.lianjia.com/chengjiao/101084782030....	101084782030	116.475489	40.019520	11110
1	https://bj.lianjia.com/chengjiao/101086012217....	101086012217	116.453917	39.881534	11110
2	https://bj.lianjia.com/chengjiao/101086041636....	101086041636	116.561978	39.877145	11110
3	https://bj.lianjia.com/chengjiao/101086406841....	101086406841	116.438010	40.076114	11110
4	https://bj.lianjia.com/chengjiao/101086920653....	101086920653	116.428392	39.886229	11110
...
160869	https://bj.lianjia.com/chengjiao/BJYZ92299606....	BJYZ92299606	116.550348	39.755625	11110
160870	https://bj.lianjia.com/chengjiao/BJYZ92303482....	BJYZ92303482	116.546899	39.755236	11110
160871	https://bj.lianjia.com/chengjiao/BJYZ92311192....	BJYZ92311192	116.546899	39.755236	11110
160872	https://bj.lianjia.com/chengjiao/BJYZ92324217....	BJYZ92324217	116.497256	39.804081	11110
160873	https://bj.lianjia.com/chengjiao/BJYZ92363553....	BJYZ92363553	116.497474	39.810115	11110

160874 rows × 26 columns



Data Preprocessing:

1. Some columns have been dropped as they aren't important in the analysis.

2. As the total null values in the dataset are only ~1%, those rows have been dropped from the dataset.
3. The column 'floor' had unrecognizable non-numeric characters along with numeric characters. The numeric characters have been extracted and the non-numeric characters deleted.
4. Some other columns that had numbers as strings have been converted to integer types for building the regression model later on.
5. The column 'tradeTime' was in string format. It was converted to a date type and then year, month and day have been extracted as new columns from it as integer types.
6. Some Variables that are categorical, their type has been defined as 'category' for building

```
In [3]: columnsToDrop = ['url', 'id', 'Cid', 'price']
df = df.drop(columns=columnsToDrop)
nullValuesPercent = round((df.isnull().sum()/len(df))*100,2)
print(f'percentage of null values by column: \n{nullValuesPercent}')
```

percentage of null values by column:

Lng	0.00
Lat	0.00
tradeTime	0.00
DOM	0.00
followers	0.00
totalPrice	0.00
square	0.00
livingRoom	0.00
drawingRoom	0.00
kitchen	0.00
bathRoom	0.00
floor	0.00
buildingType	0.82
constructionTime	0.00
renovationCondition	0.00
buildingStructure	0.00
ladderRatio	0.00
elevator	0.02
fiveYearsProperty	0.02
subway	0.02
district	0.00
communityAverage	0.13
dtype:	float64

```
In [4]: df.dropna(inplace=True)
df.shape
```

Out[4]: (159376, 22)

```
In [5]: cols = ['tradeTime', 'livingRoom', 'drawingRoom', 'bathRoom', 'floor', 'construction']  
for col in cols:  
    print(f'{col}:{df[col].unique()}')
```

```

tradeTime:['2016-08-09' '2016-07-28' '2016-12-11' ... '2014-06-02' '2014-08-
04'
'2014-05-01']
livingRoom:[2 3 1 4 5 6 0 7 '3' '1' '4' '2' '7' '5' '0' '6']
drawingRoom:[1 2 0 4 3 '1' '2' '3' '0' '4' '5']
bathRoom:[1 2 3 0 4 5 6 '2' '1' '0' '3' '4' '6' '5']
floor:['\u00c2\u00e2\u009f 26' '\u00c2\u009f 22' '\u00c2\u0096\u00c2\u0090 4' '\u00c2\u0097 21' '\u00c2\u0096\u00c2\u0090 6'
'\u00c2\u0096\u00c2\u0090 8' '\u00c2\u009f 6' '\u00c2\u009f 10' '\u00c2\u0096\u00c2\u0090 23' '\u00c2\u0097 11'
'\u00c2\u009f 24' '\u00c2\u0096\u00c2\u0090 23' '\u00c2\u0096\u00c2\u0090 19' '\u00c2\u009f 18' '\u00c2\u0096\u00c2\u0090 25'
'\u00c2\u0096\u00c2\u0090 12' '\u00c2\u0096\u00c2\u0090 14' '\u00c2\u0096\u00c2\u0090 30' '\u00c2\u0096\u00c2\u0090 27'
'\u00c2\u0096\u00c2\u0090 5' '\u00c2\u0096\u00c2\u0090 18' '\u00c2\u0097 28' '\u00c2\u0096\u00c2\u0090 11' '\u00c2\u0096\u00c2\u0090 9'
'\u00c2\u0097\u00c2\u0097 7' '\u00c2\u0097\u00c2\u0097 27' '\u00c2\u0096\u00c2\u0090 6' '\u00c2\u0096\u00c2\u0090 17' '\u00c2\u0097\u00c2\u0097 6' '\u00c2\u0096\u00c2\u0090 24'
'\u00c2\u0096\u00c2\u0090 15' '\u00c2\u0097 5' '\u00c2\u0096\u00c2\u0090 29' '\u00c2\u0097\u00c2\u0097 19' '\u00c2\u0097\u00c2\u0097 5'
'\u00c2\u0096\u00c2\u0090 9' '\u00c2\u0096\u00c2\u0090 22' '\u00c2\u0097\u00c2\u0097 18' '\u00c2\u0096\u00c2\u0090 16' '\u00c2\u009f 13'
'\u00c2\u009f 9' '\u00c2\u009f 17' '\u00c2\u0097 6' '\u00c2\u0096\u00c2\u0090 28' '\u00c2\u0096\u00c2\u0090 26'
'\u00c2\u0097 15' '\u00c2\u009f 16' '\u00c2\u0096\u00c2\u0090 7' '\u00c2\u0096\u00c2\u0090 13' '\u00c2\u0096\u00c2\u0090 33'
'\u00c2\u0097 14' '\u00c2\u009f 15' '\u00c2\u0097\u00c2\u0097 11' '\u00c2\u0096\u00c2\u0090 32' '\u00c2\u0097\u00c2\u0097 16'
'\u00c2\u0097 18' '\u00c2\u0097\u00c2\u0097 17' '\u00c2\u0096\u00c2\u0090 14' '\u00c2\u0096\u00c2\u0090 10' '\u00c2\u0097 20'
'\u00c2\u009f 12' '\u00c2\u0096\u00c2\u0090 31' '\u00c2\u0097 4' '\u00c2\u0097 2' '\u00c2\u0096\u00c2\u0090 30'
'\u00c2\u0096\u00c2\u0090 19' '\u00c2\u0096\u00c2\u0090 12' '\u00c2\u0096\u00c2\u0090 10' '\u00c2\u0096\u00c2\u0090 16' '\u00c2\u0097\u00c2\u0097 20'
'\u00c2\u0097 19' '\u00c2\u0096\u00c2\u0090 31' '\u00c2\u0096\u00c2\u0090 13' '\u00c2\u0097 10' '\u00c2\u009f 25'
'\u00c2\u0096\u00c2\u0090 21' '\u00c2\u0096\u00c2\u0090 20' '\u00c2\u009f 20' '\u00c2\u0096\u00c2\u0090 21' '\u00c2\u0096\u00c2\u0090 24'
'\u00c2\u0097\u00c2\u0097 4' '\u00c2\u009f 21' '\u00c2\u009f 7' '\u00c2\u0096\u00c2\u0090 22' '\u00c2\u0096\u00c2\u0090 7'
'\u00c2\u0097 8' '\u00c2\u0097\u00c2\u0097 15' '\u00c2\u0096\u00c2\u0090 18' '\u00c2\u009f 28' '\u00c2\u0097\u00c2\u0097 14' '\u00c2\u0097\u00c2\u0097 13'
'\u00c2\u0096\u00c2\u0090 20' '\u00c2\u0097 26' '\u00c2\u0096\u00c2\u0090 17' '\u00c2\u0097\u00c2\u0097 24' '\u00c2\u0097 23' '\u00c2\u0097\u00c2\u0097 21'
'\u00c2\u0097 24' '\u00c2\u009f 30' '\u00c2\u009f 11' '\u00c2\u0097 25' '\u00c2\u009f 27'
'\u00c2\u0097 9' '\u00c2\u0096\u00c2\u0090 11' '\u00c2\u0096\u00c2\u0090 28' '\u00c2\u0096\u00c2\u0090 15' '\u00c2\u0097\u00c2\u0097 26' '\u00c2\u0096\u00c2\u0090 34'
'\u00c2\u0097\u00c2\u0097 12' '\u00c2\u0096\u00c2\u0090 25' '\u00c2\u0097 17' '\u00c2\u009f 32' '\u00c2\u009f 8' '\u00c2\u0097\u00c2\u0097 3'
'\u00c2\u009f 19' '\u00c2\u0097 7' '\u00c2\u0097\u00c2\u0097 28' '\u00c2\u0097\u00c2\u0097 9' '\u00c2\u009f 31' '\u00c2\u0096\u00c2\u0090 26'
'\u00c2\u0097\u00c2\u0097 8' '\u00c2\u0097 32' '\u00c2\u0096\u00c2\u0090 42' '\u00c2\u0097 30' '\u00c2\u0096\u00c2\u0090 32'
'\u00c2\u0096\u00c2\u0090 3' '\u00c2\u0097 22' '\u00c2\u009f 14' '\u00c2\u009f 23' '\u00c2\u0097 27'
'\u00c2\u0097 13' '\u00c2\u0096\u00c2\u0090 27' '\u00c2\u0097\u00c2\u0097 23' '\u00c2\u0096\u00c2\u0090 29' '\u00c2\u0097\u00c2\u0097 22' '\u00c2\u009f 29'
'\u00c2\u0096\u00c2\u0090 34' '\u00c2\u0097 3' '\u00c2\u0097\u00c2\u0097 25' '\u00c2\u009f 34' '\u00c2\u0097\u00c2\u0097 10' '\u00c2\u0096\u00c2\u0090 37'
'\u00c2\u0097 16' '\u00c2\u0097 12' '\u00c2\u0096\u00c2\u0090 6' '\u00c2\u0097\u00c2\u0097 32' '\u00c2\u0096\u00c2\u0090 42'
'\u00c2\u0097\u00c2\u0097 30' '\u00c2\u0097 33' '\u00c2\u009f 42' '\u00c2\u0096\u00c2\u0090 8' '\u00c2\u009f 33'
'\u00c2\u0096\u00c2\u0090 33' '\u00c2\u0097\u00c2\u0097 2' '\u00c2\u0097\u00c2\u0097 29' '\u00c2\u0097 1' '\u00c2\u0097 29'
'\u00c2\u0096\u00c2\u0090 15' '\u00c2\u009f 37' '\u00c2\u0096\u00c2\u0090 36' '\u00c2\u0096\u00c2\u0090 35' '\u00c2\u0097\u00c2\u0097 34'
'\u00c2\u009f 36' '\u00c2\u0096\u00c2\u0090 37' '\u00c2\u0096\u00c2\u0090 35' '\u00c2\u0097 31'
'\u00c2\u0096\u00c2\u0090 12' '\u00c2\u0097\u00c2\u0097 31' '\u00c2\u0096\u00c2\u0090 63' '\u00c2\u0096\u00c2\u0090 21'
'\u00c2\u0097 34' '\u00c2\u0096\u00c2\u0090 57' '\u00c2\u0097\u00c2\u0097 33' '\u00c2\u0096\u00c2\u0090 11'
'\u00c2\u0096\u00c2\u0090 10' '\u00c2\u0096\u00c2\u0090 8' '\u00c2\u0096\u00c2\u0090 18'
'\u00c2\u0096\u00c2\u0090 7' '\u00c2\u0096\u00c2\u0090 20' '\u00c2\u0096\u00c2\u0090 25'
'\u00c2\u0096\u00c2\u0090 16' '\u00c2\u0096\u00c2\u0090 23' '\u00c2\u0096\u00c2\u0090 14'
'\u00c2\u0096\u00c2\u0090 27' '\u00c2\u0096\u00c2\u0090 28' '\u00c2\u0096\u00c2\u0090 5'
'\u00c2\u0096\u00c2\u0090 22' '\u00c2\u0096\u00c2\u0090 17' '\u00c2\u0096\u00c2\u0090 2' '\u00c2\u0096\u00c2\u0090 4'
'\u00c2\u0096\u00c2\u0090 13' '\u00c2\u0096\u00c2\u0090 36' '\u00c2\u0096\u00c2\u0090 26'
'\u00c2\u0096\u00c2\u0090 24' '\u00c2\u0096\u00c2\u0090 9' '\u00c2\u0096\u00c2\u0090 30'
'\u00c2\u0096\u00c2\u0090 19' '\u00c2\u0096\u00c2\u0090 31']
constructionTime:['2005' '2004' '2008' '1960' '1997' '2009' '1991' '2001' '1
990' '2011'
'2000' '1998' '2010' '1996' '1993' '2006' '2002' '\u00c2\u0096\u00c2\u0090 2012'
'1989' '2003' '2007' '1994' '1984' '1992' '2014' '1985' '1999' '1979'
'1981' '1976' '1982' '1975' '1983' '1986' '1995' '1965' '2013' '1988'
'1987' '2015' '1955' '1980' '1978' '1958' '1970' '1956' '1977' '1964'
'1963' '1967' '2016' '1974' '1973' '1959' '1954' '1962' '1966' '1957'
'1972' '1971' '1953' '1968' '1961' '1950' '1952' '1969']

```

```
In [6]: def extractNumeric(s):
        return re.sub(r'\D', '', s)
df['floor'] = df['floor'].apply(extractNumeric)

cols = cols[1:]
for col in cols:
    df[col] = pd.to_numeric(df[col], errors='coerce')

meanConstructionTime = round(df['constructionTime'].mean())
df['constructionTime'].fillna(meanConstructionTime, inplace=True)

df['tradeTime'] = pd.to_datetime(df['tradeTime'])
df['tradeYear'] = df['tradeTime'].dt.year
df['tradeMonth'] = df['tradeTime'].dt.month
df['tradeDay'] = df['tradeTime'].dt.day
df.drop(columns=['tradeTime'], inplace=True)
```

```
In [7]: catCols = ['buildingType', 'renovationCondition', 'buildingStructure', 'elevator']
for col in catCols:
    df[col] = df[col].astype('category')
```

Regression Model:

Light Gradient Boosting Machine (LGBM) is chosen for building the regression model. LGBM is chosen because it is known for its efficiency in training and inference, especially on large datasets. It implements gradient boosting algorithms, which generally train faster than traditional ensemble methods like Random Forest. It also generally outperforms both Random Forest and Linear Regression in terms of predictive accuracy.

The column 'totalPrice' is chosen as the target variable. The dataset is split into training and testing sets. Based on the training results the top 10 important features out of the 23 features are selected. Testing is done once with all the features and then with only the top 10 features. The RMSE in both cases is same ~107. Also the R² value in both the cases is ~81%.

```
In [8]: # Dividing the data into target variable and the independent variables and also
X = df.drop(columns=['totalPrice'])
y = df['totalPrice']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [9]: # Defining the parameters of the Lgbm model
params = {'boosting_type': 'dart',
          'objective': 'regression',
          'metric': 'rmse',
          'num_leaves': 31,
          'learning_rate': 0.05,
          'feature_fraction': 0.9,
          'seed': 99}

trainData = lgb.Dataset(X_train, label=y_train, categorical_feature=catCols)
testData = lgb.Dataset(X_test, label=y_test, categorical_feature=catCols, ref=
```

```
In [10]: model = lgb.train(params, trainData, num_boost_round=100, valid_sets=[testData])
```

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.009210 seconds.

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1904

[LightGBM] [Info] Number of data points in the train set: 127500, number of used features: 23

[LightGBM] [Info] Start training from score 409.377336

```
In [11]: # Important Variables in predicting the price in decreasing order
featureImportance = model.feature_importance(importance_type='gain')

featureImportanceDf = pd.DataFrame({'Feature': model.feature_name(), 'Importance': featureImportance})

featureImportanceDf = featureImportanceDf.sort_values(by='Importance', ascending=False)

print(featureImportanceDf)
```

	Feature	Importance
4	square	6.622094e+10
19	communityAverage	5.334922e+10
20	tradeYear	1.739085e+10
2	DOM	1.407755e+10
18	district	5.806078e+09
8	bathRoom	1.796832e+09
21	tradeMonth	7.694184e+08
5	livingRoom	5.460086e+08
11	constructionTime	1.325027e+08
0	Lng	1.298598e+08
1	Lat	1.247739e+08
14	ladderRatio	7.813243e+07
3	followers	4.208859e+07
15	elevator	1.716025e+07
12	renovationCondition	8.658620e+06
13	buildingStructure	3.162620e+06
10	buildingType	0.000000e+00
16	fiveYearsProperty	0.000000e+00
17	subway	0.000000e+00
9	floor	0.000000e+00
7	kitchen	0.000000e+00
6	drawingRoom	0.000000e+00
22	tradeDay	0.000000e+00

```
In [12]: # Predicting on the test set using all the features
y_pred = model.predict(X_test, num_iteration=model.best_iteration)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))

r2 = r2_score(y_test, y_pred)

print("RMSE:", rmse)
print("R^2:", r2)
```

```
RMSE: 107.88755984222088
R^2: 0.8144671341774744
```

```
In [13]: topFeatures = featureImportanceDf.head(10)['Feature'].tolist()
print(f'The 10 most important features are: \n{topFeatures}')
```

```
The 10 most important features are:
['square', 'communityAverage', 'tradeYear', 'DOM', 'district', 'bathRoom',
'tradeMonth', 'livingRoom', 'constructionTime', 'Lng']
```



```
In [14]: # Predicting on the test set using only the top 10 features
newCatCol = ['district']

X_train_top = X_train[topFeatures]
X_test_top = X_test[topFeatures]

trainDataTop = lgb.Dataset(X_train_top, label=y_train, categorical_feature=newCatCol)
testDataTop = lgb.Dataset(X_test_top, label=y_test, categorical_feature=newCatCol)

modelTop = lgb.train(params, trainDataTop, num_boost_round=100, valid_sets=[testDataTop])
y_pred_top = modelTop.predict(X_test_top, num_iteration=modelTop.best_iteration)
```

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001517 seconds.

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1136

[LightGBM] [Info] Number of data points in the train set: 127500, number of used features: 10

[LightGBM] [Info] Start training from score 409.377336

```
In [15]: rmseTop = np.sqrt(mean_squared_error(y_test, y_pred_top))

r2Top = r2_score(y_test, y_pred_top)

print("RMSE:", rmseTop)
print("R^2:", r2Top)
```


RMSE: 107.75417981563106

R^2: 0.8149255944818011

Below is a scatter plot showing most of the predicted prices lying within a 95% confidence interval of the actual prices.

```
In [16]: confidence_interval = 1.96 * np.std(y_pred_top)

plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_top, alpha=0.5)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], '--', color=
plt.plot([min(y_test), max(y_test)], [min(y_test) - confidence_interval, max(y
plt.plot([min(y_test), max(y_test)], [min(y_test) + confidence_interval, max(y
plt.title('Actual Price vs Predicted Price with Confidence Interval')
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.legend()
plt.show()
```



Classification of Properties as Underpriced or Overpriced.

1. A subset of the dataframe consisting of only the top ten features and the target variable 'totalPrice' is created. This is to reduce the curse of dimensionality while using K-Means. Even though the variable 'Lat': Latitude is not in the top 10 features, it is included because it is used later on in the analysis to locate the properties on a map.
2. As the computational speed becomes very slow with a large dataset while using K-Means, so a random sample of 2000 rows is selected to work on.
3. Since the dataset consists of a mix of continuous and categorical variables, it is preprocessed by calculating the Gower's distance of the variables before plugging into the K-Means model.

```
In [17]: subsetDf = df[topFeatures]
subsetDf['totalPrice'] = df['totalPrice']
subsetDf['Lat'] = df['Lat']
```

C:\Users\himan\AppData\Local\Temp\ipykernel_20144\4027769544.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
subsetDf['totalPrice'] = df['totalPrice']
```

C:\Users\himan\AppData\Local\Temp\ipykernel_20144\4027769544.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
subsetDf['Lat'] = df['Lat']
```

```
In [18]: # Taking a random sample of 2000
randomSample = subsetDf.sample(n=2000, random_state=78, replace=False)
selectedCols = [col for col in randomSample.columns if col != 'totalPrice']
selectedData = randomSample[selectedCols]

selectedData = selectedData.apply(lambda x: x.astype('category').cat.codes if

# Calculating the gower's distance
distanceMatrix = gower_matrix(selectedData)
```


Number of Clusters:

The ideal number of clusters to be used in K-Means can be decided by different methods. Two methods are used here:

1. The Elbow Method: The idea is to select the value of k at the "elbow" point, as it represents the point where increasing the number of clusters does not significantly decrease the within-cluster sum of squares or the Inertia. Based on the plot the cluster value at the elbow is 3.
2. The Silhouette Score Method: The objective of clustering is to group similar data points together while keeping dissimilar points in different clusters. A high average Silhouette Score indicates that this objective is being achieved effectively. Based on this method the ideal cluster value is 2.


```
In [19]: inertia = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(distanceMatrix)
    inertia.append(kmeans.inertia_)

plt.plot(range(1, 11), inertia, marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Inertia')
plt.title('Elbow Method')
plt.show()
```



```
In [20]: silhouetteScores = []
for k in range(2, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(distanceMatrix)
    score = silhouette_score(distanceMatrix, kmeans.labels_)
    silhouetteScores.append(score)

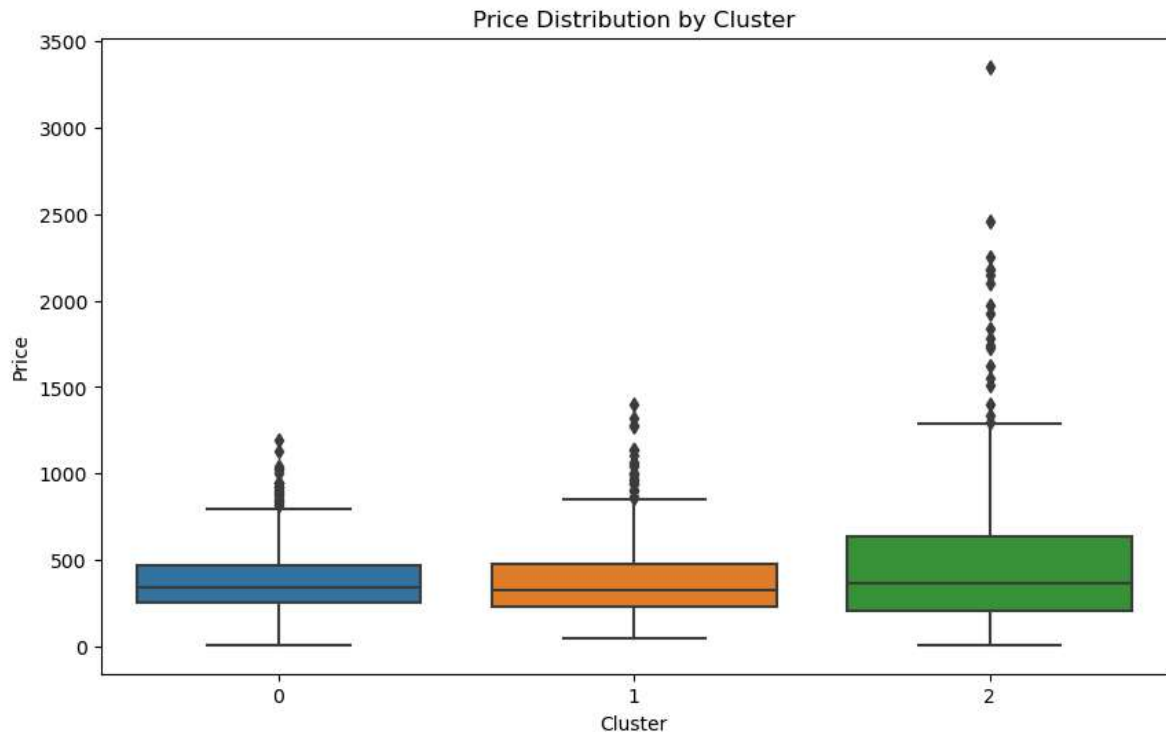
plt.plot(range(2, 11), silhouetteScores, marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Silhouette Score')
plt.title('Silhouette Score Method')
plt.show()
```



```
In [21]: # no. of clusters = 3
kmeans = KMeans(n_clusters=3, n_init=10)
kmeans.fit(distanceMatrix)
cluster_labels = kmeans.labels_
```

The properties are divided among 3 clusters based on the 10 features selected earlier. The following graph and the summary statistics are based on the prices of properties within each cluster.

```
In [22]: clusterPrices = pd.DataFrame({'Cluster': cluster_labels, 'Price': randomSample
plt.figure(figsize=(10, 6))
sns.boxplot(x='Cluster', y='Price', data=clusterPrices)
plt.title('Price Distribution by Cluster')
plt.xlabel('Cluster')
plt.ylabel('Price')
plt.show()
```



```
In [23]: clusterSummaryStats = clusterPrices.groupby('Cluster')['Price'].describe()

print("Summary statistics for each cluster:")
print(clusterSummaryStats)
```

Summary statistics for each cluster:

	count	mean	std	min	25%	50%	75%	max
Cluster								
0	952.0	381.227521	176.919966	11.2	253.00	345.0	472.0	1190.0
1	627.0	380.086284	213.558944	50.0	229.75	330.0	480.0	1400.0
2	421.0	498.329454	428.490219	11.1	205.00	368.0	638.0	3350.0

The method of classifying a property as underpriced or overpriced is based on calculating the mean price and standard deviation of each cluster. Any property within a cluster, if its price is more than one standard deviation than the mean, it is classified as overpriced. If its price is less than one standard deviation of the mean price of its cluster, then it is classified as underpriced.

The following table shows the datapoints classified as such. Of the 2000 data points analyzed, 483 or around 1/4th are either underpriced or overpriced.

```

In [24]: clusterAvgPrice = clusterPrices.groupby('Cluster')['Price'].mean()

clusterPrices['ClusterAvgPrice'] = clusterPrices['Cluster'].map(clusterAvgPrice)

clusterStd = clusterPrices.groupby('Cluster')['Price'].std()

clusterPrices['ClusterStd'] = clusterPrices['Cluster'].map(clusterStd)

thresholdMultiplier = 1
overpriced = clusterPrices[clusterPrices['Price'] > clusterPrices['ClusterAvgPrice'] * thresholdMultiplier]
underpriced = clusterPrices[clusterPrices['Price'] < clusterPrices['ClusterAvgPrice'] * thresholdMultiplier]

overpricedIndices = overpriced.index
underpricedIndices = underpriced.index

overpricedDf = randomSample.loc[overpricedIndices].copy()
underpricedDf = randomSample.loc[underpricedIndices].copy()
overpricedDf = overpricedDf.assign(PricedType='Overpriced')
underpricedDf = underpricedDf.assign(PricedType='Underpriced')

overOrUnderPricedDf = pd.concat([overpricedDf, underpricedDf])

overOrUnderPricedDf

```

Out[24]:

	square	communityAverage	tradeYear	DOM	district	bathRoom	tradeMonth	livingRoom
37042	51.01	119449.0	2016	16.0	10	1	9	
79302	107.30	60013.0	2017	68.0	7	1	8	
60670	107.41	94271.0	2017	43.0	1	1	3	
35429	83.00	75858.0	2016	5.0	8	1	9	
88863	117.47	52197.0	2017	20.0	6	2	11	
...
14148	54.00	46072.0	2016	25.0	7	1	7	
95227	61.89	37661.0	2015	1.0	6	1	3	
118720	37.93	59730.0	2016	1.0	7	1	5	
93145	128.71	38281.0	2011	1.0	6	2	12	
19158	35.50	58959.0	2016	57.0	2	1	9	

483 rows × 13 columns

The following is a street map of Beijing showing the overpriced (in Red) and underpriced (in Green) properties.

```
In [25]: map = folium.Map(location=[overOrUnderPricedDf ['Lat'].mean(), overOrUnderPricedDf ['Lng'].mean()], overOrUnderPricedDf,
        colors = {'Underpriced': 'green', 'Overpriced': 'red'})

for index, row in overOrUnderPricedDf .iterrows():
    folium.Marker([row['Lat'], row['Lng']], icon=folium.Icon(color=colors[row['Status']]))

map
```

Out[25]:

